Zen AI Model Family

Zen-Nano

Mobile/IoT Intelligence

Technical Whitepaper v1.0

Hanzo AI Research Team
research@hanzo.ai

Zoo Labs Foundation
foundation@zoolabs.org

September 2025

Abstract

We present **Zen-Nano**, a 0.6B parameter model optimized for mobile/iot intelligence. Built upon zen-0.5B, this model achieves state-of-the-art performance while maintaining exceptional efficiency with only 0.6B active parameters. Supporting 64K thinking tokens for advanced reasoning, the model represents a significant advancement in democratizing AI through sustainable and efficient architectures.

Contents

1 Introduction

The rapid advancement of artificial intelligence has created an unprecedented demand for models that balance capability with efficiency. **Zen-Nano** addresses this challenge by delivering enterprise-grade performance while maintaining a minimal computational footprint.

1.1 Key Innovations

- **Efficient Architecture**: 0.6B active parameters from 0.6B total

- **Specialized Training**: Optimized for mobile/iot intelligence

- **Extended Context**: 32K context window

- **Thinking Mode**: 64K thinking tokens

2 Architecture

2.1 Model Design

Zen-Nano is based on the zen-0.5B architecture with several key modifications:

Component	Specification
Total Parameters	0.6B
Active Parameters	0.6B
Base Model	zen-0.5B
Context Length	32K
Thinking Tokens	64K
Architecture Type	Transformer

Table 1: Zen-Nano Architecture Specifications

2.2 Technical Innovations

2.2.1 Mixture of Experts (MoE)

The model uses a dense architecture with all parameters active during inference, optimized for maximum performance per parameter.

2.2.2 Attention Mechanism

Extended attention mechanisms support up to 32K context length with efficient KV-cache management.

2.2.3 Thinking Mode

Advanced reasoning through extended thinking tokens (up to 64K), enabling:

- Step-by-step problem decomposition

- Self-correction and verification

- Complex multi-step reasoning

- Internal deliberation before response

3 Performance Benchmarks

3.1 Evaluation Results

Benchmark	Score
MMLU	51.7%
HumanEval	22.6%
GSM8K	62.0%
HellaSwag	59.5%

Table 2: Language Understanding Benchmarks

3.2 Efficiency Metrics

Metric	Value
Inference Speed	450 tokens/sec
Memory Usage (INT4)	2 GB
Energy Efficiency	98% reduction
Latency (First Token)	15 ms

Table 3: Efficiency Metrics

4 Training Methodology

4.1 Dataset

The model was trained on a carefully curated dataset comprising:

- High-quality filtered web data (0.5TB)

- Domain-specific corpora for mobile/iot intelligence

- Synthetic data generation for edge cases

- Human feedback through RLHF

4.2 Training Process

1. **Pretraining**: 2 trillion tokens over 14 days on 8x A100

2. **Supervised Fine-tuning**: Task-specific optimization

3. **RLHF**: Alignment with human preferences

4. **Constitutional AI**: Safety and helpfulness optimization

5 Use Cases and Applications

5.1 Primary Applications

Conversational AI and chatbots

Content generation and summarization

Code completion and review

Educational assistance

Research and analysis

5.2 Integration Examples

```
from transformers import AutoModelForCausalLM , AutoTokenizer


# Load model and tokenizer
model = AutoModelForCausalLM.from_pretrained("zenlm/zen-nano-0.6b-
    instruct")
tokenizer = AutoTokenizer.from_pretrained("zenlm/zen-nano-0.6b-instruct
    ")


# Generate response
inputs = tokenizer("Explain␣quantum␣computing", return_tensors="pt")
outputs = model.generate(**inputs , max_length =100)
response = tokenizer.decode(outputs [0])
```

Listing 1: Basic Usage Example

6 Environmental Impact

6.1 Sustainability Metrics

- **Carbon Footprint**: 0.02 kg COe per million inferences

- **Energy Usage**: 0.5 kWh per day (1000 users)

- **Efficiency Gain**: 98% reduction vs comparable models

6.2 Green AI Commitment

Zen AI models are designed with sustainability as a core principle, achieving industry-leading efficiency through architectural innovations and optimization techniques.

7 Safety and Alignment

7.1 Safety Measures

- Constitutional AI training for harmlessness

- Comprehensive red-teaming and adversarial testing

- Built-in safety filters and guardrails

- Regular safety audits and updates

7.2 Ethical Considerations

The model has been developed with careful attention to:

- Bias mitigation through diverse training data

- Transparency in capabilities and limitations

- Privacy-preserving deployment options

- Responsible AI principles alignment

8 Deployment Options

8.1 Available Formats

- **SafeTensors**: Original precision weights

- **GGUF**: Quantized formats (Q4_K_M, Q5_K_M, Q8_0)

- **MLX**: Apple Silicon optimization (4-bit, 8-bit)

- **ONNX**: Cross-platform deployment (coming soon)

8.2 Hardware Requirements

Precision	Memory	Recommended Hardware
FP16	1.2 GB	RTX 3060
INT8	0.6 GB	GTX 1660
INT4	2 GB	Raspberry Pi 5

Table 4: Hardware Requirements by Precision

9 Future Work

9.1 Planned Improvements

- Extended context windows (up to 1M tokens)

- Enhanced multimodal capabilities

- Improved efficiency through further optimization

- Expanded language support

9.2 Research Directions

- Advanced reasoning mechanisms

- Self-supervised learning improvements

- Zero-shot generalization enhancement

- Continual learning capabilities

10 Conclusion

Zen-Nano represents a significant advancement in AI democratization, delivering exceptional performance for mobile/iot intelligence while maintaining unprecedented efficiency. Through innovative architecture design and careful optimization, the model achieves a balance between capability and sustainability that sets a new standard for responsible AI development.

Acknowledgments

References

A   Model Card

Field	Value
Model Name	Zen-Nano
Version	1.0.0
Release Date	September 2025
License	Apache 2.0
Repository	huggingface.co/zenlm/zen-nano-0.6b-instruct
Documentation	github.com/zenlm/zen
Contact	research@hanzo.ai

Table 5: Model Card Information