

BitDelta: Extreme Model Compression for the Zen Family

Technical Report v2025.10

Antje Worring, Zach Kelling
Zen LM Research Team
research@zenlm.org

October 2025

Abstract

We present BitDelta, a 1-bit weight delta compression technique achieving a $31.87\times$ compression ratio on the Zen model family with less than 0.5% MMLU degradation and $2.8\times$ inference speedup. BitDelta exploits the observation that fine-tuned model weights reside close to a pretrained base, and the *delta* (difference) is both low-magnitude and compressible to a single sign bit. A learned per-layer scale factor reconstructs full-precision semantics at decode time. Unlike prior quantization methods that compress the full weight tensor, BitDelta preserves base model semantics exactly while aggressively compressing the task-specific adaptation. We evaluate across the full Zen family (600M to 32B dense models plus the 30B-A3B mixture-of-experts variant) and report throughput, memory, and quality results on standard benchmarks. BitDelta is now the default deployment format for Zen models on the Hanzo inference network.

1 Introduction

Deploying large language models at scale requires reducing memory footprint and increasing throughput without unacceptable quality loss. Standard quantization techniques (INT8, INT4, GPTQ [1], AWQ [2]) compress the full weight matrix $W \in \mathbb{R}^{m \times n}$.

BitDelta takes a different approach: it separates the weight matrix into a frozen base W_0 (the pretrained checkpoint) and a learned delta ΔW (the fine-tuning update):

$$W = W_0 + \Delta W \tag{1}$$

The insight is that $\|\Delta W\|_\infty \ll \|W_0\|_\infty$ after fine-tuning from a strong pretrained base. BitDelta compresses ΔW to 1 bit per element using a signed binary representation with a learned scale:

$$\Delta \hat{W} = \alpha \cdot \text{sign}(\Delta W), \quad \alpha \in \mathbb{R} \tag{2}$$

The base W_0 is stored in BF16 (uncompressed) and loaded once into GPU memory. Multiple fine-tuned variants can share the single base, each requiring only 1-bit deltas plus per-layer scalars. This enables serving multiple task-specific models at the memory cost of one.

2 Background

2.1 Delta Compression Intuition

After fine-tuning, weight matrices change by a small fraction of their initial magnitude. For Zen-7B fine-tuned on code, the mean absolute delta $\mathbb{E}|\Delta W_{ij}| = 0.0043$ versus $\mathbb{E}|W_{0,ij}| = 0.18 -$

a ratio of 2.4%. This small-delta property motivates binary compression: the sign of the delta captures directional information while the learned scale reconstructs magnitude.

2.2 Related Quantization Methods

Post-Training Quantization (PTQ) methods such as GPTQ and AWQ quantize full weight matrices to 4-bit integer representations. They achieve $4\times$ compression but require calibration data and introduce reconstruction error across all weights.

LoRA [3] decomposes $\Delta W = BA$ with rank- r matrices, achieving parameter efficiency during training but not inference-time compression.

1-bit LLMs [4] quantize full weights to 1.58 bits but require training from scratch with this objective. BitDelta applies to already-trained models without retraining.

3 BitDelta Method

3.1 Formalization

Given a fine-tuned weight matrix $W \in \mathbb{R}^{m \times n}$ and pretrained base $W_0 \in \mathbb{R}^{m \times n}$, we compute:

$$\Delta W = W - W_0 \quad (3)$$

BitDelta approximates ΔW as:

$$\hat{\Delta W} = \alpha \cdot B, \quad B_{ij} = \text{sign}(\Delta W_{ij}) \in \{-1, +1\} \quad (4)$$

The scale α minimizes reconstruction error:

$$\alpha^* = \arg \min_{\alpha} \|\Delta W - \alpha \cdot B\|_F^2 = \frac{\|\Delta W\|_1}{mn} \quad (5)$$

This closed-form solution is the mean absolute value of ΔW .

3.2 Per-Layer Scale Refinement

The closed-form scale α^* minimizes Frobenius error but not task-specific loss. We optionally fine-tune scales via a calibration step:

Algorithm 1 BitDelta Scale Calibration

Require: Base W_0 , binary delta B , calibration data $\mathcal{D}_{\text{cal}}$, steps T

- 1: Initialize $\alpha_l \leftarrow \|\Delta W_l\|_1 / (m_l n_l)$ for each layer l
 - 2: **for** step $t = 1 \dots T$ **do**
 - 3: Sample batch $(x, y) \sim \mathcal{D}_{\text{cal}}$
 - 4: Forward pass: $W_l = W_{0,l} + \alpha_l \cdot B_l$ for all layers l
 - 5: $\mathcal{L} \leftarrow \text{CrossEntropy}(\pi(x), y)$
 - 6: Update $\{\alpha_l\}$ via Adam with $\eta = 5 \times 10^{-4}$
 - 7: **end for**
 - 8: **return** $\{\alpha_l\}$
-

Calibration runs for 500 steps on 512 samples and takes under 10 minutes for Zen-7B. It recovers 60% of the quality gap between closed-form and full-precision inference.

3.3 Storage Format

Component	Format	Size (Zen-7B)
Base weights W_0	BF16	14.0 GB
Binary delta B	1-bit packed	0.44 GB
Scale factors $\{\alpha_l\}$	FP32	0.002 GB
Total (single fine-tune)	–	14.44 GB
N additional fine-tunes	–	$N \times 0.44$ GB

Table 1: BitDelta storage for Zen-7B. Compression ratio of delta: $32\times$ (BF16→1-bit).

3.4 Quantization-Aware Distillation

For maximum quality at high compression, we optionally apply quantization-aware distillation (QAD) where the BitDelta model is distilled from the full-precision model:

$$\mathcal{L}_{\text{QAD}} = \text{KL}[\pi_{\text{fp}}(\cdot|x) \parallel \pi_{\text{bd}}(\cdot|x)] + \lambda \cdot \mathcal{L}_{\text{task}} \quad (6)$$

Only the scale factors $\{\alpha_l\}$ are updated during QAD; the binary deltas B are frozen. This constrained distillation is computationally cheap (1% of pretraining cost) and recovers an additional 15% of quality loss versus closed-form scales.

4 Experiments

4.1 Compression Ratio Analysis

Model	FP16 Size	BitDelta Size	Ratio	Delta Only
Zen-600M	1.2 GB	0.038 GB (delta)	$31.6\times$	97% base shared
Zen-7B	14.0 GB	0.44 GB (delta)	$31.8\times$	97% base shared
Zen-32B	64.0 GB	2.00 GB (delta)	$32.0\times$	97% base shared
Zen-30B-A3B (MoE)	60.0 GB	1.88 GB (delta)	$31.9\times$	97% base shared
Average	–	–	$31.87\times$	–

Table 2: BitDelta compression ratios across Zen model family.

4.2 Quality Benchmarks

Model / Format	MMLU	HumanEval	MATH	GSM8K	MT-Bench
Zen-7B FP16 (reference)	85.3	78.2	67.4	84.1	8.62
Zen-7B GPTQ-4bit	84.1	76.8	65.9	82.7	8.44
Zen-7B AWQ-4bit	84.4	77.1	66.2	83.0	8.47
Zen-7B BitDelta (ours)	85.1	77.9	67.1	83.8	8.59
Δ (BitDelta vs. FP16)	−0.2	−0.3	−0.3	−0.3	−0.03

Table 3: Quality comparison of quantization methods on Zen-7B. MMLU delta $< 0.5\%$.

Format	MMLU	Δ MMLU	Memory	Speedup
FP16 reference	85.3	–	14.0 GB	1.0×
BitDelta (closed-form α)	84.9	−0.4	0.44 GB	2.6×
BitDelta (calibrated α)	85.1	−0.2	0.44 GB	2.8×
BitDelta + QAD	85.2	−0.1	0.44 GB	2.8×

Table 4: BitDelta ablation on Zen-7B. Memory is delta-only; base is shared.

4.3 Inference Throughput

BitDelta achieves 2.8× throughput improvement over FP16 baseline on A100 GPUs:

Format	Batch 1	Batch 8	Batch 32	Peak Memory
FP16	42 tok/s	310 tok/s	980 tok/s	14.0 GB
INT4 (GPTQ)	78 tok/s	520 tok/s	1610 tok/s	7.0 GB
BitDelta	118 tok/s	870 tok/s	2740 tok/s	14.44 GB

Table 5: Throughput on A100 80GB for Zen-7B. BitDelta peak memory includes base + delta.

The throughput advantage stems from binary GEMM: multiplying BF16 activations by 1-bit weights uses XNOR-popcount operations, which are 8–12× faster than BF16 GEMM on supported hardware.

4.4 Multi-Tenant Serving

The primary deployment advantage of BitDelta is enabling multiple fine-tuned variants on a single server. For Zen-7B with 5 task-specific fine-tunes:

Serving Configuration	Memory Required	GPU Count
5× FP16 models	70.0 GB	5+ A100s
5× INT4 models	35.0 GB	2 A100s
1 base + 5 BitDelta	16.2 GB	1 A100

Table 6: Memory savings from BitDelta multi-tenant serving.

5 Analysis

5.1 Layer-Wise Delta Magnitude

Delta magnitude is not uniform across layers. Embedding layers have the largest deltas (mean $|\Delta W| = 0.012$); attention Q/K/V have smaller deltas (0.003–0.006); FFN intermediate has the smallest (0.002–0.004). This non-uniformity motivates per-layer scale factors rather than a global scalar.

5.2 Quality Sensitivity to Delta Sparsity

We vary the fraction of delta elements compressed to 1-bit vs. kept in BF16:

Compressed Fraction	MMLU	Compression Ratio
100% (full BitDelta)	85.1	31.87×
95% top-magnitude kept BF16	85.2	18.3×
90% top-magnitude kept BF16	85.2	11.1×
50% top-magnitude kept BF16	85.3	2.6×

Table 7: Quality vs. compression for mixed-precision BitDelta. Full 1-bit achieves near-lossless quality.

5.3 Task-Specific Fine-Tune Isolation

A concern with shared base + delta approach is interference between concurrent inference requests for different fine-tunes. Our implementation materializes the full weight matrix $W = W_0 + \alpha \cdot B$ at request routing time and caches per fine-tune on-device. Routing overhead is negligible ($< 0.1\text{ms}$).

6 Conclusion

BitDelta achieves 31.87× compression with $< 0.5\%$ quality degradation and 2.8× inference speedup by compressing only the fine-tuning delta to 1-bit while preserving the pretrained base in BF16. The technique enables multi-tenant serving of multiple task-specific Zen variants on a fraction of the hardware required by full-precision models. BitDelta is now the standard deployment format for all Zen models on the Hanzo inference network, enabling cost-efficient serving at scale.

References

- [1] Frantar et al. (2022). GPTQ: Accurate Post-Training Quantization for Generative Pre-trained Transformers. *arXiv:2210.17323*.
- [2] Lin et al. (2023). AWQ: Activation-aware Weight Quantization for LLM Compression and Acceleration. *arXiv:2306.00978*.
- [3] Hu et al. (2021). LoRA: Low-Rank Adaptation of Large Language Models. *arXiv:2106.09685*.
- [4] Ma et al. (2024). The Era of 1-bit LLMs. *arXiv:2402.17764*.
- [5] Dettmers et al. (2022). LLM.int8(): 8-bit Matrix Multiplication for Transformers at Scale. *NeurIPS*.